


IDENTIFYING SARCASM ON TWITTER

Rachel M. Adler



PROBLEM DEFINITION

- Businesses and government organizations often need to be able to identify sarcasm online (e.g., the Secret Service¹)
- Identifying sarcasm is also vital to accurate sentiment analysis²
- However, determining whether a message is sarcastic is difficult for both people and computers³

¹ Hannon et al., 2004 ² Filatova, 2012; Maynard & Greenwood, 2014 ³ Wallace et al., 2014

GOAL

Develop a model using Twitter data to predict whether a message is sarcastic

THE DATA

- Data collected via Twitter REST API between July and October, 2017
- Gathered most recent 3,200 tweets made by each of John Oliver's 3,573,510 followers
- Used tweets with 4 hashtags to build models (#sarcasm, #happy, #sad, #seriously)
- Use sarcasm hashtag (*#sarcasm*) as true label


THE DATA

Summary of tweets acquired

| | Total Tweets | Unique Users |
|------------|--------------|--------------|
| #sarcasm | 30,910 | 23,509 |
| #happy | 12,639 | 10,149 |
| #sad | 40,861 | 26,456 |
| #seriously | 11,450 | 9,033 |

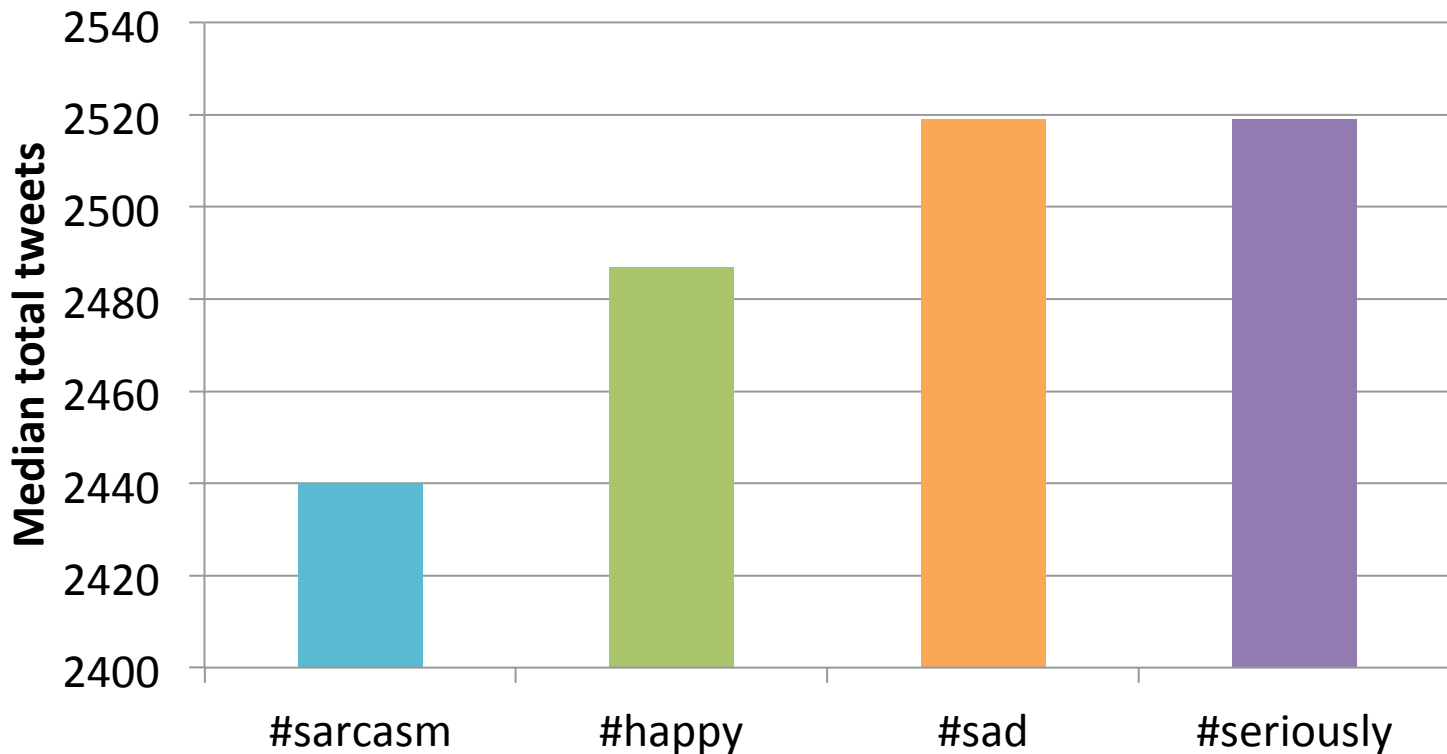
EXPLORATORY DATA ANALYSIS

Median number of followers per user



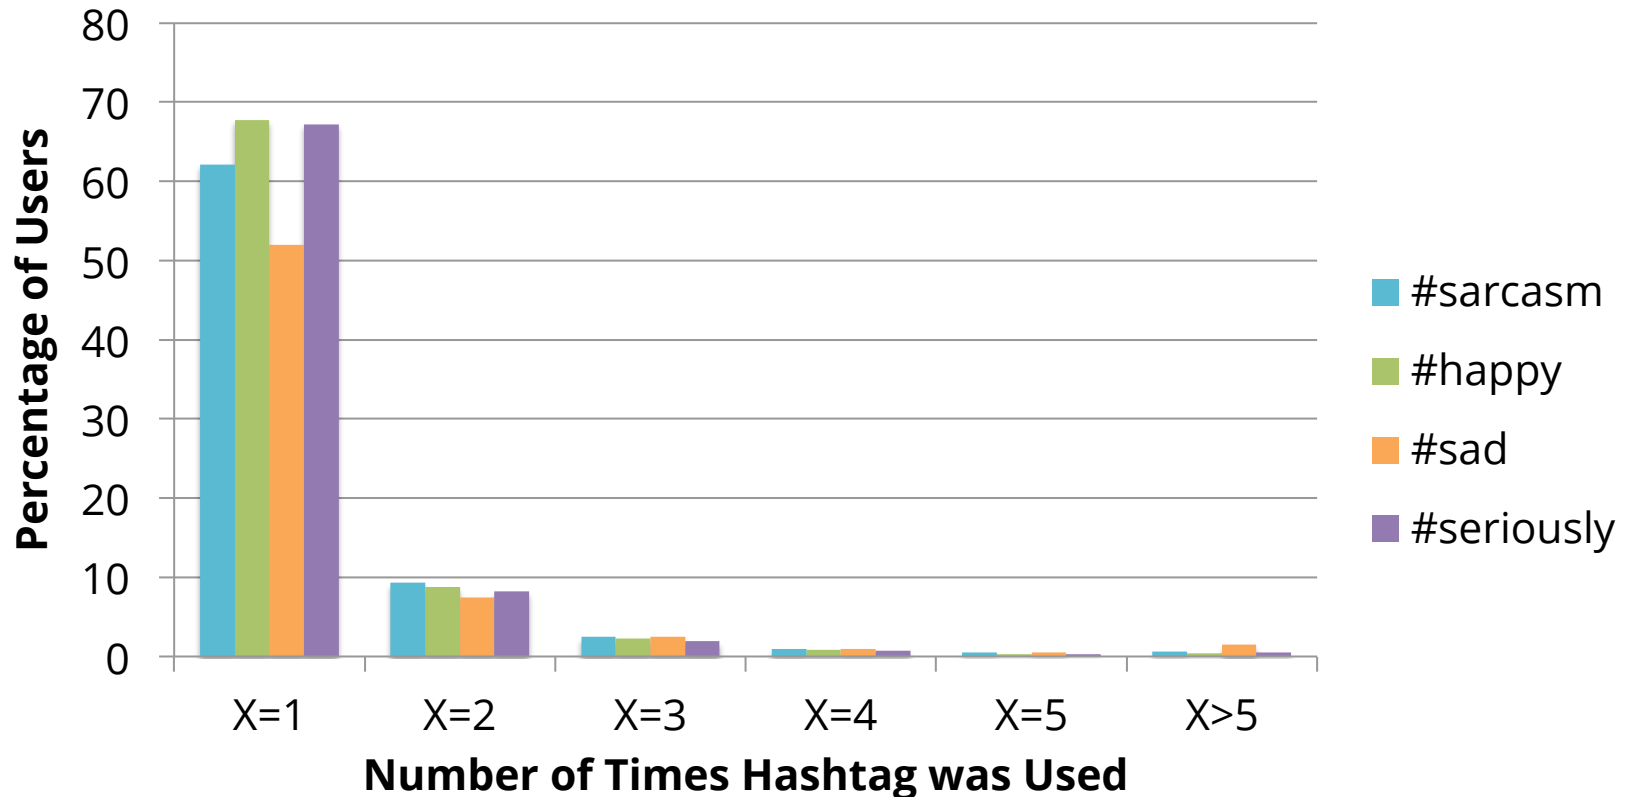
EXPLORATORY DATA ANALYSIS

Median total tweets per user



EXPLORATORY DATA ANALYSIS

Percentage of users that used each hashtag X times



MODELING SETUP

- Selected subset of data for modeling so that tweets with #sarcasm represented about 50% of dataset

| | Total Tweets | Percentage of Tweets |
|------------|--------------|----------------------|
| #sarcasm | 30,728 | 50.14% |
| #happy | 10,040 | 16.38% |
| #sad | 10,271 | 16.76% |
| #seriously | 10,262 | 16.74% |

- Divided data into training and test sets (70-30% split) with approximately equal proportions of #sarcasm in each

BASELINE MODEL

- Created term-document matrix to represent tweets
- Generated bag of words representation, removing stop words, with scikit-learn's *CountVectorizer*
- Performed 5-fold cross-validation with AUC as scoring function
- Used Naïve Bayes classifier; see results below

| Model | Feature Vectors | AUC | Precision | Recall | F1 Score | Accuracy (Train) | Accuracy (Test) |
|-------------|------------------|------|-----------|--------|----------|------------------|-----------------|
| Naïve Bayes | Term occurrences | 0.59 | 0.53 | 0.89 | 0.66 | 57.05% | 54.12% |

NAÏVE BAYES MODEL

- Created term-document matrix to represent tweets
- Transformed matrix to TF-IDF representation
- Generated bag of words representation, removing stop words, with scikit-learn's *TfidfVectorizer*
- Performed 5-fold cross-validation with AUC as scoring function
- Used Naïve Bayes classifier; see results below

| Model | Feature Vectors | AUC | Precision | Recall | F1 Score | Accuracy (Train) | Accuracy (Test) |
|-------------|------------------|------|-----------|--------|----------|------------------|-----------------|
| Naïve Bayes | Term frequencies | 0.57 | 0.52 | 0.92 | 0.66 | 54.97% | 53.41% |

LATENT SEMANTIC INDEXING

- Used LSI to perform dimensionality reduction and to generate topics as follows:
 1. Created vocabulary for corpus
 - Excluded words appearing in fewer than 10 or more than 40% of tweets
 2. Generated bag of words representation with gensim's *doc2bow*
 3. Created three LSI models with gensim's *LsiModel*
 - 200 components
 - 300 components
 - 400 components

LSI MODELS

- Tested 3 classifiers
 - Random Forest
 - XGB Classifier
 - Logistic regression
- Tested each classifier with three LSI models (200, 300, and 400 components)
- For each model, performed 5-fold cross-validation with AUC as scoring function

RANDOM FOREST MODELS

- Table below summarizes model performance
- While AUC is higher than baseline model (AUC=0.59), high training accuracy suggests over-fitting

| Model | Feature Vectors | AUC | Precision | Recall | F1 Score | Accuracy (Train) | Accuracy (Test) |
|---------------|--------------------|------|-----------|--------|----------|------------------|-----------------|
| Random Forest | 200 LSI components | 0.67 | 0.63 | 0.60 | 0.62 | 99.70% | 62.57% |
| Random Forest | 300 LSI components | 0.65 | 0.62 | 0.57 | 0.59 | 99.71% | 60.80% |
| Random Forest | 400 LSI components | 0.65 | 0.61 | 0.57 | 0.59 | 99.71% | 60.78% |

XGB CLASSIFIER MODELS

- Table below summarizes model performance
- Models with 200 and 300 components are worse than baseline model (AUC=0.59), though model performs better with 400 components

| Model | Feature Vectors | AUC | Precision | Recall | F1 Score | Accuracy (Train) | Accuracy (Test) |
|----------------|--------------------|------|-----------|--------|----------|------------------|-----------------|
| XGB Classifier | 200 LSI components | 0.55 | 0.53 | 0.88 | 0.66 | 55.97% | 54.81% |
| XGB Classifier | 300 LSI components | 0.57 | 0.54 | 0.72 | 0.62 | 57.24% | 55.75% |
| XGB Classifier | 400 LSI components | 0.61 | 0.57 | 0.67 | 0.61 | 59.62% | 57.86% |

LOGISTIC REGRESSION MODELS

- Table below summarizes model performance
- All models outperform baseline model
- Model with 400 components performs best

| Model | Feature Vectors | AUC | Precision | Recall | F1 Score | Accuracy (Train) | Accuracy (Test) |
|---------------------|--------------------|------|-----------|--------|----------|------------------|-----------------|
| Logistic Regression | 200 LSI components | 0.72 | 0.67 | 0.63 | 0.65 | 66.15% | 66.95% |
| Logistic Regression | 300 LSI components | 0.74 | 0.68 | 0.65 | 0.66 | 69.13% | 67.33% |
| Logistic Regression | 400 LSI components | 0.74 | 0.68 | 0.65 | 0.67 | 70.04% | 67.70% |

RECOMMENDATIONS

Overall, the logistic regression model with 400 LSI components performed best and should be used for classifying sarcastic tweets

FUTURE WORK

- To make model more generalizable:
 - Expand negative examples in dataset to tweets beyond those using #happy, #sad, or #seriously
 - Use a more randomly-selected sample of tweets (rather than just tweets made by John Oliver's followers)
- Adjust proportion of positive and negative examples so there are much fewer positive examples (to better reflect distribution of tweets in real world)